



**BAJAJ
FINSERV**

INDIA'S MOST
DIVERSIFIED
FINANCIAL
SERVICES GROUP

BAJAJ
ATOM
CASE COMPETITION
9TH EDITION

FINANCIAL PROPOSAL

Bajaj Finserv Motor Insurance (BFDL)

Driving Growth. Protecting India.



PROTECT

Comprehensive protection
you can trust



GROW

Accelerating profitable
and sustainable growth



DELIGHT

Delivering seamless
customer experiences



CREATE VALUE

Generating long-term value
for all stakeholders



INNOVATE

Leveraging technology
and data for impact



PARTNER

Building strong partnerships
for shared success



FOCUS

Customer-centric approach
with sharp execution



COMPLY

Upholding integrity and
regulatory excellence



SUSTAIN

Driving responsible business
for a better tomorrow

AUGUST 2026



BAJAJ FINSERV
MOTOR INSURANCE
BFDL

CONTENTS

Note 1	The Unpriced Asset — Bajaj's Historical Claim Media	4
Note 2	Pre-Investment Measurement Plan.....	5
Note 3	Industry, Customer and Regulatory Shift.....	6
Note 4	The ClaimPulse Solution — Parallel Orchestration.....	8
Note 5	Stakeholder Map and Anticipated Challenges.....	9
Note 6	Scenario Simulation	11
Note 7	Landscape Impact — Where the Durable Advantage Relocates	13
Note 8	Resilience and Breakpoints.....	14
Note 9	Basis of Preparation.....	16
Note 10	Key Assumptions and Sources	17
Note 11	Revenue and Benefit Derivation.....	19
Note 12	Build Cost Schedule	22
Note 13	Annual Operating Cost Schedule	24
Note 14	AI Cost Architecture — The Nine Optimisation Decisions	25
Note 15	Inference Cost Sensitivity and Breakeven.....	28
Note 16	What Optimisation Is Worth.....	30
Note 17	GenAI Build versus Buy.....	31
Note 18	Net Result and Payback	33
Note 19	Insurance P&L Bridge — Combined Ratio Impact.....	34
Note 20	Downside and Correlated Stress Tests	35
Note 21	Stakeholder Value Split.....	37
Note 22	Break-Even Analysis.....	38
Note 23	Claim Frequency Reconciliation.....	39
Note 24	Alternative Net Earned Premium Bases.....	40
Note 25	Three-Year Forecast by Implementation Plan	41
Note 26	Unit Economics per Claim.....	42
Note 27	Project DCF, NPV and IRR.....	43
Note 28	ROIC and Economic Value Added	44
Note 29	Capital Allocation — Build versus Buy versus Wait.....	45
Note 30	The Cost of Waiting.....	46
Note 31	Strategic Moat and Early Warning System	48
Note 32	Implementation Roadmap.....	50
Note 33	The Investment Ask	51

NOTE 1 THE UNPRICED ASSET — BAJAJ'S HISTORICAL CLAIM MEDIA

The model values Bajaj's historical claim-image archive at zero. Training-data line D-13 carries Rs 0.60 Cr to acquire or build a labelled damage and synthetic-media set, as though ClaimPulse were starting from nothing — despite Bajaj General having settled motor claims with photographic evidence for two decades.

Item	Value	Basis
Motor OD claims per year, full addressable book	245,820	Model output at 100% rollout
Policies in force	48.3 million	Tier 1 Filed, Q1 FY27 disclosure
Retained claim images/video, historical	Not measured	Bajaj claims operations only
Images per claim	Not measured	Depends on capture policy over time

Why it matters, if it exists

#	Effect	Detail
1	D-13 falls	Rs 0.60 Cr of build cost is buying training data; an existing labelled corpus reduces or removes this line entirely.
2	The fraud graph starts warm	A graph built on years of settled claims, garages, workshops and claimants can detect rings from day one, rather than needing history to accumulate first.
3	It is the commoditisation answer	When component AI commoditises, the moat has to move to data and ecosystem. This archive is that moat, and no competitor can buy it.

DPDP purpose limitation is a real, unresolved obstacle: media collected to settle a claim was not necessarily collected with consent to train models on it. The DPDP Rules 2025 compliance date of 13 May 2027 falls inside this investment's horizon, and both the retention basis and the training basis need a legal opinion before any rupee value is claimed. This is why this note carries no rupee figure of its own — it is stated as a strategic option, not booked benefit.

NOTE 2 PRE-INVESTMENT MEASUREMENT PLAN


Six specific measurements convert this model's remaining Tier 4 assumptions into Tier 1 measured inputs. None of them changes the investment decision on its own — Notes 15, 16 and 17 already show the case is insensitive to most of them — but each closes a named gap before the figure it feeds is quoted without qualification.

#	Measurement	Why It Blocks a Number	Who Provides It
1	GPU seconds per claim, by lane	Drives 46% of run cost; every figure in Note 16 moves with it.	Pilot instrumentation
2	Tokens per claim, by pipeline stage	Rate cards are verified; volumes are not. Note 15 collapses to a single column once this exists.	countTokens over a pilot sample
3	Video seconds actually needed for damage assessment	Sets the achievable reduction in Decision 4 (Note 14.4), the dominant cost term.	Pilot, against surveyor ground truth
4	Achievable peak-to-average ratio under queueing	The 2.5x multiplier on the largest variable line (Decision 6, Note 14.6).	Load pattern from live claim arrival data
5	Retained claim media volume and retention period	Note 1 cannot be sized without this; D-13 may be overstated.	Bajaj claims operations
6	Legal basis for training on historical claim media	Determines whether Note 1 is an asset or a prohibited use.	External counsel, DPDP scope

These six measurements form the evidence-conversion mechanism behind the phased capital release described in the main investment proposal: each pilot gate exists specifically to close one of these gaps before the next tranche of capital is committed.

NOTE 3 INDUSTRY, CUSTOMER AND REGULATORY SHIFT

Every figure in this document sits inside an external environment this case does not control. Twenty-one named trends across four domains — Consumer, Technology, Industry, Regulatory — were tracked, each carrying a named publisher, a date, and an evidence class. Eighteen are verified clean against a live source and quotable without qualification; nine are carried at a lower tier and none of those is load-bearing on a rupee figure, except the graph-AI detection-uplift range, which only bounds a target this model already stress-tests to zero (Note 30).



The chart plots each trend's expected magnitude against how well-evidenced it is. Governance requirements (T9) and explainability expectation (C5) sit in the highest-confidence, highest-relevance corner — both are legal facts, not projections, so the architecture treats them as production requirements rather than compliance overlays (Note 14). Synthetic-media capability (T5) and fraud behaviour (S3) sit at the lowest evidence confidence despite high expected magnitude: no published incidence of synthetic claim media exists anywhere in the Indian motor market, a gap independently rediscovered by all four research engines. This is the single data point that would most strengthen this case if it could be obtained.

3.1 The shift, by domain

Domain	What is shifting	Evidence	Implication for ClaimPulse
Consumer	Digital-first settlement is now the default expectation, carried over from UPI and digital KYC rather than from insurance itself	Mordor Intelligence: mobile app 56.1%, P&C 41.8% of India's online insurance market	Speed is now a baseline expectation, not a differentiator — the moat has to sit elsewhere (Note 5)
Consumer	Only a minority are comfortable with AI deciding a claim unaided; support for AI assistance is broader	Insurity: 20-39% support AI assistance; only 22% comfortable with AI filing a claim unaided	Human escalation must stay visible and always available — never a fully autonomous product

Domain	What is shifting	Evidence	Implication for ClaimPulse
Technology	Generative image tools have made photographic evidence cheap to fabricate	NVIDIA/ControlExpert case study: 95% detection accuracy on manipulated motor-claim images	The Capture Integrity Gate is a first-class build line (Note 12), not a later addition
Industry	General insurance is in a soft pricing cycle; underwriting results are stressed market-wide	Bajaj Finserv Q1 FY27: combined ratio 104.7%, up from 103.6% a year earlier	This is a current-quarter fact driving the case, not a forecast (Note 19)
Regulatory	IRDAI moved insurers to proactive, board-governed fraud monitoring	IRDAI Fraud Monitoring Framework Guidelines, effective 1 Apr 2026	Fraud governance is a compliance obligation from day one, not a competitive choice
Regulatory	Motor losses below Rs 50,000 no longer require a registered surveyor; app-based AI assessment is explicitly contemplated	IRDAI Master Circular, 11 Jun 2024	The Green-lane design sits inside an exemption IRDAI itself created, not a grey area
Regulatory	National data-protection obligations land inside this investment's five-year horizon	DPDP Rules 2025: full compliance 13 May 2027, penalties to Rs 250 Cr, no grace period	Data residency and consent architecture are costed build lines (Note 12), not legal risk carried unpriced

3.2 What the trend work settled

Seven design decisions trace directly to a named trend rather than to team preference: the evidence-integrity layer became a costed build line because T-04 makes every downstream model unsafe without it; the fraud target was capped and stress-tested because the graph-AI evidence range (T-05) sits above published benchmarks; repair-network gains were kept out of the investment case because T-07 delivers customer benefit that does not accrue to Bajaj directly; rollout was made the primary lever rather than a fixed plan because T-09 means deployment must track the pricing decision, not lead it; and regulatory tightening is priced as a build overlay (Note 17, Scenario S5) rather than ignored, because T-13's eventual framework is a known unknown.

NOTE 4 THE CLAIMPULSE SOLUTION — PARALLEL ORCHESTRATION

ClaimPulse replaces the sequential, seven-touch claims process with five AI engines that score every claim concurrently against the same data, fusing their outputs into a single Trust Score that routes the claim into a Green, Amber or Red lane. This is an operating-model change, not a faster version of the existing sequence: the delay in the current process comes from queue-stacking, where each function waits for the previous one to finish before it can start, and concurrency removes that mathematical source of delay directly rather than speeding up each queue individually.

4.1 Why parallel, not cascaded

A cascaded design — cheap checks first, expensive checks only if needed — was considered and rejected. It reintroduces exactly the queue-stacking this architecture exists to remove, and low-value claims are a known adversarial-structuring target, so short-circuiting scrutiny on them is a fraud-surface risk, not a genuine cost saving. Every claim is scored by all five engines before it is ever routed; this concurrency is the non-negotiable architectural constraint behind the TAT reduction quoted throughout this document.

4.2 The five engines

Engine	What it decides	Build cost ref.
Capture Integrity Gate	Whether submitted evidence is genuine, unmanipulated and internally consistent	D-03 to D-05
Document AI / Policy RAG	What the claim documents and the actual policy wording say, extracted and cross-validated	D-01, D-02
CV Damage Assessment	Damage severity, location and an indicative repair-cost band	D-06, D-07
Fraud and Duplicate Graph	Whether this claim connects to other claims, claimants, garages or devices in a way that signals organised fraud	D-08
Trust Score Orchestration	Fuses all four upstream outputs into one score and routes the claim to Green, Amber or Red	D-09

The governing rule behind automation eligibility is stated explicitly in the underlying model: digitalisation times evidence integrity gives automation eligibility, not digital evidence volume on its own. A claim can be fully digital and still unsafe to auto-settle if the evidence behind it is not verifiably genuine — which is precisely why the Capture Integrity Gate sits ahead of, not alongside, the other four engines in build priority (Note 12).

NOTE 5 STAKEHOLDER MAP AND ANTICIPATED CHALLENGES

Note 21 quantifies what each stakeholder receives in rupees. This note carries the qualitative side: what each stakeholder needs today, what changes for them, and the specific objection each is likely to raise, stated in their own terms rather than smoothed into a general FAQ.

5.1 Stakeholder map

Stakeholder	What they need today	What ClaimPulse changes
Policyholder / claimant	Certainty and speed — status apps are typically abandoned after 2-3 silent visits	Green-lane auto-settlement in about 1.5 days instead of 9.8, with proactive status push rather than pull
Bajaj General claims operations	Fewer low-value manual touches, and capacity for genuinely complex claims	7 manual touches to 1.23 blended; adjusters move from data entry to exception handling
Underwriting and the CFO	Loss-ratio and expense-ratio relief without cutting distribution	Fraud leakage down and claims opex down — both sides of the combined ratio move together
BFDL / Bajaj Markets	A differentiated distribution proposition and a platform asset the group owns	Owns the Trust Score engine and Capture Integrity Gate as a group capability, extensible to health and home
Bajaj Finserv group leadership	Capital allocated to something that compounds, not a one-off saving	A proprietary fraud graph and synthetic-media defence that improve with every claim processed
Surveyors	Deployment against claims that actually need judgement	AI triage handles sub-Rs 50,000 claims inside the IRDAI exemption; surveyors focus above it
Garages and the repair network	Faster approvals and predictable estimates	The repair-cost engine gives an indicative band at first notification from settled claims and live parts feeds
IRDAI and the regulator	Explainable, auditable, fair automated decisions	Human override, full audit trail and per-decision explanation built into production hardening
Honest claimants who cannot use live capture	Not to be penalised for an older phone, indoor parking, or GPS switched off	The Amber lane absorbs them with an innocent explanation rather than a rejection — the friction is costed explicitly, not hidden (Note 21)


5.2 What each stakeholder will challenge, and the answer

Who asks	The challenge	Our answer
Claims operations	Seven manual touches is not our number.	Agreed — it is a team estimate, flagged Tier 4, and the model recalculates the moment the real figure is provided.
CFO / finance	A three-month payback is not credible.	On steady-state economics it is three months. From project kickoff, including the ten-month build, it is about seventeen (Note 27). If both weakest assumptions break at once, it is still under nine.
CFO / finance	Rs 250 per touch is outsourced pricing; our staff are in-house.	Correct — the downside case reprices every touch at Rs 83, and the case still clears (Note 20).
Underwriting	You cannot subtract a Motor OD saving from our group combined ratio.	We do not. Both bases are shown separately with the denominator stated on every line (Note 19).
Actuarial	Your fraud leakage number comes from health insurance.	No motor-specific Indian leakage figure is published. The proxy is stated with its basis, and closing it is the top data request (Note 2).
Data science	Ninety percent fraud precision is above the published range for graph models.	It is a target, not a validated capability, and it is labelled that way throughout. The downside case lands it at 82% (Note 20).
Group finance	This is under half a percent of group PAT.	On Motor OD alone. Motor is the proof, not the prize — the gate and the graph extend to health and home without re-architecture (Note 31).
Legal and compliance	Live-capture-only will exclude some honest claimants.	It will, and we cost it: 8% of claims fall to Amber, a named Rs 0.59 Cr annual cost at Base. They are still handled, just not auto-settled.

Who asks	The challenge	Our answer
Strategy	What stops a competitor copying this in eighteen months?	The orchestration is copyable. The fraud graph and synthetic-media classifier are trained on Bajaj's own claim media and compound with every claim processed — that is the moat (Note 7).
Anyone	Why is BFDL funding a benefit that lands in Bajaj General's P&L?	At group level, Bajaj Finserv consolidates both — a capital allocation question, not a transfer question. An explicit internal service arrangement is still recommended.

NOTE 6 SCENARIO SIMULATION

Every scenario below moves one lever and holds everything else at Base, so its effect is attributable to that lever alone — Level 2 of Scenario S1 is the Base case by construction, which is how the engine checks itself. All figures tie to the live model on six independent anchors before being quoted.



Scenario	Level	Lever value	Net annual, Rs Cr	Delta vs Base
S1 Digital-first settlement	L1 Emerging	55% green	42.10	-1.21
S1 Digital-first settlement	L2 Established (= Base)	65% green	43.31	0
S1 Digital-first settlement	L3 Accelerated	78% green	44.88	+1.57
S2 Synthetic evidence	L1 Gate holds	90% detection	43.31	0
S2 Synthetic evidence	L2 Partial erosion	77% detection	32.41	-10.90
S2 Synthetic evidence	L3 Detection neutralised	62% detection	19.82	-23.49
S5 Regulatory AI governance	L1 No new rule	+0% build	43.31	0
S5 Regulatory AI governance	L2 EU AI Act lower bound	+15% build	43.31	0 (payback +1.0 mo)
S5 Regulatory AI governance	L3 EU AI Act upper bound	+20% build	43.31	0 (payback +1.4 mo)

Scenario S3 (explainability requirement) carries no rupee figure, since no credible source prices an explainability obligation for Indian motor claims; a Rs 2.81 Cr/year grievance-cost exposure is shown as an order of magnitude only. Scenario S4 (repair-ecosystem digitisation) improves customer-facing TAT from 2.55 to 1.76 days but changes ClaimPulse's own economics by nothing, and is deliberately kept out of the investment case for that reason.

6.1 Compound stress test — all three quantified threats at once

Lever	Base	Compound severe
Auto-settlement share	65%	50%
Fraud detection achieved	90%	62% (full revert to rule-based)
Build cost, Rs Cr	9.89	11.87
NET ANNUAL BENEFIT, Rs Cr	43.31	15.95
PAYBACK, steady state, months	2.74	8.90

Lever	Base	Compound severe
COMBINED RATIO, Motor OD, pp	2.58	1.04
<i>There is no combination of the researched scenarios in which ClaimPulse fails to return its capital inside a year of steady-state operation. The one genuine failure mode sits outside this table: Note 8 documents the single case, across every scenario tested by every research engine, where monthly economics turn negative.</i>		

NOTE 7 LANDSCAPE IMPACT — WHERE THE DURABLE ADVANTAGE RELOCATES

Seventeen scenarios were tested across four independent research engines (Customer, Technology, Social, Industrial). Two conclusions recur across all of them and are stated here as the headline findings, ahead of the detail.

7.1 Two headline findings

Finding	What it means
Speed is not the moat.	TAT-based differentiation compresses industry-wide as every insurer chases the same benchmark. All four research engines independently relocate the durable advantage to the same four places: proprietary claims data, garage and ecosystem relationships, governance built ahead of the regulatory mandate, and the ability to retrain faster than a competitor. The scenario that looks financially benign — component commoditisation — is the one that threatens the moat most, because it removes technology as a differentiator and leaves only data and relationships standing.
Regulatory tightening is not a competitive threat.	Three separate engines reach the same finding: because an IRDAI constraint applies to every insurer simultaneously, ClaimPulse's relative position is preserved even as its absolute automation ceiling falls. Regulation compresses payback (Note 6, Scenario S5); it does not compress the value proposition.

7.2 Three landscape themes, synthesised

Theme	Mechanism	What it does to the competitive landscape
Experience parity and digital visibility	Customers benchmark claims against UPI and food delivery, not against other insurers — a reference point no single insurer can opt out of resetting.	Weakens TAT-as-differentiator industry-wide; strengthens ClaimPulse's relative position only if the evidence-integrity layer (Note 4) is also built, since speed alone stops being defensible once the whole industry converges on it.
Governed autonomy and AI regulation	The automation frontier is set by permission, not capability — governance is what buys the right to automate, not a compliance afterthought.	Converts the competitive question from 'who has the best AI' to 'who can prove the best AI is safely governed.' A platform with governance embedded from the start is harder for a fast-follower to replicate than one with governance retrofitted after a regulatory shock.
Adversarial claims and fraud evolution	Detection capability and generation capability both rise; the variable that matters is the gap between them, not either in absolute terms.	The one scenario that genuinely threatens the case (Note 6, S2 Level 3) — and even there, the labour engine alone still repays the build within about a year.

7.3 Competitive convergence — bounded, not point-estimated


A rival matching the full ClaimPulse pipeline is the single largest expected-loss line in the underlying risk model, at 50% probability and roughly 15% benefit erosion from year two — a 7.5% expected loss. It is graded qualitatively rather than run as a financial scenario, because erosion depends on a rival's rollout pace and pricing response, neither of which is observable today. It is bounded instead: the fraud-detection scenarios in Note 6 already test benefit erosion of 25% and 54% from a different cause, and a 15% competitive erosion sits comfortably inside the range this model is shown to survive.

NOTE 8 RESILIENCE AND BREAKPOINTS

Do not say ClaimPulse is resilient. Say where it breaks. Nine breakpoints are defined below; five are computable from the live model today, four are not. The fraud floor has already been breached in testing and the model survived it: with fraud benefit set entirely to zero, the case still returns Rs 3.20 Cr a year and repays the build in 37 months — the honest floor, not the three-month headline.

8.1 The nine breakpoints

Ref	Breakpoint	Status
STP*	Straight-through processing floor — the Green-lane share at which NPV turns zero	Computable. Break-even rollout on labour savings alone is 36.19% at Base (Note 22)
FD*	Fraud-detection floor — the capability level at which fraud economics stop contributing	Computable and demonstrated (chart below). Absolute floor: net Rs 3.20 Cr/year, 37-month payback
C*	Complexity ceiling — the Amber/Red migration share at which operating economics break	Requires lane calibration from the pilot
H*	Human exception-capacity ceiling — the exception volume at which P90 TAT starts to deteriorate	Requires Bajaj claims-operations data. Flagged by two engines as the constraint most likely to bind first — the pilot's first measurement (Note 2)
Alcost*	AI/inference cost ceiling	Computable from the live model (Notes 15-16)
V*	Addressable volume floor for positive NPV	Computable from the live model
G*	Governance cost ceiling before ROI falls below the hurdle rate	Partially computable — Note 6 S5 tests +15% and +20% of build
AG*	Adversarial gap floor — the point at which generation capability overtakes detection	Bounded, not computed: rests on synthetic-media generation capability, for which no published series exists
OO*	Orchestration-ownership threshold — the point an external party captures the orchestration layer	Qualitative only; no quantification source exists



8.2 Three resilience dimensions — reported separately, never averaged

Dimension	Definition	Status
Economic resilience	NPV under the scenario divided by NPV at Base	Computable per regime
Operational resilience	Deterioration in P90 TAT, human workload, STP, reinspection and exception rate	Requires a P90 distribution, which the pilot produces

Dimension	Definition	Status
Architectural resilience	Whether the architecture can absorb new fraud vectors, models, data sources or governance requirements without rebuilding the core	Assessable today, qualitative

These three are deliberately never collapsed into one weighted score. A composite resilience index is defined in the underlying model (a weighted sum of model, operational, financial and strategic resilience) but its weights are a business-judgement input, not a research output, so no single number is presented here without the weights that produced it stated alongside it.

8.3 The one failure mode found across every engine and every scenario

Across all four research engines and every grade and level tested, exactly one case turns net monthly benefit negative: a combined technology stress where automation depth, human capacity and fraud detection collapse simultaneously (roughly minus Rs 0.53 Cr/month at the extreme grade). The model does not fail under any single-mechanism shock, however severe — it fails only when three things break together. That is the honest boundary of this case, stated rather than left to be found under questioning.

NOTE 9 BASIS OF PREPARATION

This model is built bottom-up from Bajaj's own filed disclosures wherever a filed figure exists, supplemented by vendor-published rate cards and a small number of team-estimated operating assumptions, each labelled by evidence tier (Tier 1 Filed, Tier 1 Vendor, Tier 2, Tier 4, Target, or Standard). Sheets 2 to 13 of the source workbook contain no typed numbers other than the assumption inputs on Sheet 1 — every other cell is a formula. This annexure preserves that discipline: every figure below traces to a named reference cell in the R5 model.

Construction conventions

Item	Convention Applied	Where It Lands
FX rate	Spot, Rs 95.3 / USD (week to 11 Aug 2026 range: 94.89–95.45)	Every USD-benchmarked build line
Build cost	Bottom-up, fifteen costed components	Note 12
Run cost	Bottom-up, six lines (two variable, four fixed)	Note 13
Fraud benefit	Leakage pool x reduction in the undetected share	Note 11, Part C
Surveyor exemption threshold	Rs 50,000 (IRDAI Master Circular, 2024)	Lane design
Capture Integrity Gate	Costed in build lines D-03 to D-05, D-13; valued in benefit lines W-20, W-21	Notes 11, 12
TAT	Lane-weighted, restated after live-capture friction	Note 11, Part B

Three implementation plans — Conservative (20% rollout), Base (60%), Aggressive (100%) — are carried side by side throughout. The only two levers that differ between them are the rollout percentage and the post-implementation cost per manual touch (which falls with volume through BPO discounts); build cost, all rate cards, and all per-claim assumptions are identical across the three.

NOTE 10 KEY ASSUMPTIONS AND SOURCES

10.1 Filed and external data

Ref	Item	Value	Tier
A-01	Motor OD GDPI, FY2025-26	Rs 3,317.3 Cr	Tier 1 Filed
A-02	Net earned premium / GDPI	49.88%	Tier 1 Filed
A-03	Net incurred claims ratio	74.6%	Tier 1 Filed
A-04	Group combined ratio, FY2024-25	102.3%	Tier 1 Filed
A-05	Bajaj total GDPI, FY2025-26	Rs 20,461 Cr	Tier 1 Filed
A-06	Motor OD market growth p.a.	8.7%	Tier 2
A-07	Salary inflation p.a.	6.9%	Tier 1
A-08	FX rate	Rs 95.3 / USD	Market
A-09	Discount rate (WACC)	12.0%	Standard
A-10	India regulated-sector cost uplift	1.314x	Tier 2
A-11	IRDAI surveyor exemption threshold	Rs 50,000	Tier 1
A-12	API / inference price deflation p.a.	-30%	Tier 2
A-13	GPU price deflation p.a.	-10%	Tier 2

Table A. Source: GI Council segment filings, Bajaj General FY2026 and Q1 FY2027 results, IRDAI Financial Highlights, spot FX August 2026.

10.2 Operating assumptions

Ref	Assumption	Value	Tier
B-01	Average claim size	Rs 45,000	Tier 4
B-02	Manual touches per claim, today	7	Tier 4
B-03 / 04 / 05	Lane mix — Green / Amber / Red	65% / 25% / 10%	Tier 2/4
B-06 / 07 / 08	Touches — Green / Amber / Red	0.2 / 2 / 6	Tier 4
B-09	Claim TAT today	9.8 days	Tier 2
B-10 / 11 / 12	TAT — Green / Amber / Red	1.5 / 3.5 / 7.0 days	Tier 4
B-13	Fraud detection rate, today	62%	Tier 4
B-14	Fraud detection rate, target	90%	Target
B-15	Uplift share — Capture Integrity Gate	10 pp of 28	Tier 4
B-16	Uplift share — fraud/duplicate graph	18 pp of 28	Tier 4
B-17	Fraud leakage pool, % of claims value	3.15%	Tier 2 Proxy
B-18	Synthetic-media fraud, % of claims value	1.00%	Tier 4
B-19	Gate detection rate, synthetic media	85%	Tier 4
B-20	Live-capture friction (claims downgraded)	8%	Tier 4
B-21	Average Motor OD premium per policy	Rs 7,000	Tier 4
B-22	Renewal uplift on claimants	5 pp	Tier 4
B-23	Retained value on premium	10%	Tier 4
B-24 / 25 / 26	Benefit realisation — Yr 1 / Yr 2 / Yr 3	45% / 92% / 100%	Plan
B-27	Build duration	10 months	Plan
B-28	Risk haircut on net benefit	23.75%	Tier 4
B-28	Cost per manual touch, baseline	Rs 250	Tier 4

Table B. B-14 (90% target) is tested at an 82% downside landing (F-02) throughout this document; every place the 90% figure is used, the stress-tested figure is shown alongside it.

10.3 The three implementation plans


Ref	Lever	Conservative	Base	Aggressive
C-01	Rollout of the addressable Motor OD book	20%	60%	100%
C-02	Cost per manual touch, post-implementation	Rs 300	Rs 250	Rs 200

Table C. These two levers are the only differences between the plans. Cost per touch falls as volume rises through BPO volume discounts and applies only to post-implementation touches; the baseline rate (B-28) is fixed. Build cost is identical across all three, because the same system is built regardless of rollout.

NOTE 11 REVENUE AND BENEFIT DERIVATION

11.1 From filed premium to claims in scope

Ref	Step	Conservative	Base	Aggressive	Unit
W-01	Motor OD GDPI, FY2025-26	3,317.30	3,317.30	3,317.30	Rs Cr
W-02	Motor OD net earned premium	1,654.67	1,654.67	1,654.67	Rs Cr
W-03	Annual incurred claims pool	1,234.38	1,234.38	1,234.38	Rs Cr
W-04	Average claim size	45,000	45,000	45,000	Rs
W-05	Annual claim count, full book	274,307	274,307	274,307	claims
W-06	Rollout of the addressable book	20%	60%	100%	ratio
W-07	Annual claims on the platform	54,861	164,584	274,307	claims
W-08	Monthly claims on the platform	4,572	13,715	22,859	claims
W-09	Annual claims value on the platform	246.88	740.63	1,234.38	Rs Cr




The claims pool is built entirely from filed premium: Motor OD GDPI (Tier 1 Filed) times the filed net-earned-premium ratio times the filed net-incurred-claims ratio gives the annual claims pool in rupees, which is then divided by the average claim size to give a claim count. Rollout scales the count and value that reach the platform; nothing else in this table is a free variable.

11.2 Lane mix, manual touches and TAT — the Trust Score routing

Ref	Step	Conservative	Base	Aggressive	Unit
W-10	Blended manual touches after ClaimPulse	1.23	1.23	1.23	touches
W-11	Manual touches removed per claim	5.77	5.77	5.77	touches
W-12	Claim TAT on ClaimPulse	2.55	2.55	2.55	days
W-13	Claim TAT today	9.80	9.80	9.80	days
W-14	TAT on ClaimPulse, after live-capture friction	2.71	2.71	2.71	days
W-15	Blended TAT across the whole Motor OD book	8.38	5.55	2.71	days
W-16	TAT reduction vs today	1.42	4.25	7.09	days
W-17	TAT reduction	14.5%	43.4%	72.3%	%


Trust Score Routing — Lane-Weighted Touches and TAT



Blended touches (1.23) and blended TAT (2.55 days) are lane-weighted at 65% Green x 0.2 touches / 1.5 days, 25% Amber x 2 touches / 3.5 days, and 10% Red x 6 touches / 7 days — identical across all three plans, because lane mix is a routing-design constant, not a rollout variable. The whole-book blended TAT (W-15) is the only row that moves with rollout, since it weights platform TAT against the unreached 9.8-day baseline.

11.3 Annual benefit lines

Ref	Benefit line	Conservative	Base	Aggressive	Unit
W-18	Labour: manual touches removed	7.58	23.74	41.26	Rs Cr
W-19	Fraud avoided: fraud/duplicate graph	3.68	11.05	18.42	Rs Cr
W-20	Fraud avoided: Capture Integrity Gate	2.05	6.14	10.23	Rs Cr
W-21	Synthetic-media exposure avoided	2.10	6.30	10.49	Rs Cr
W-22	Renewal retention, claimant cohort	0.19	0.58	0.96	Rs Cr
W-23	Less: live-capture friction cost	-0.24	-0.59	-0.79	Rs Cr
W-24	GROSS ANNUAL BENEFIT	15.36	47.21	80.57	Rs Cr



Labour benefit is claims processed times the touches removed, priced at the plan's touch rate — it lands in the expense ratio. Both fraud lines apply the leakage pool to the undetected share of claims value, split 18pp/28 to the graph engine and 10pp/28 to the Capture Integrity Gate; both land in the loss ratio. Synthetic-media exposure avoided is a new,

additive vector (claims value x 1% incidence x 85% gate detection) that nothing else in the industry currently defends against. Live-capture friction is subtracted, not omitted: 8% of otherwise-clean claims fall from Green to Amber because the claimant could not complete live capture, and that cost is charged against the benefit before the gross figure is struck.


Every rupee in this section is anchored to a filed premium figure or a vendor-published rate. The only Tier-4 (team-estimated) multipliers are the touch count, lane mix, and leakage-pool proxy, each logged with its validation route in the sensitivity register (Table H/N, Sheet 1).

NOTE 12 BUILD COST SCHEDULE

The build is a one-off cost, identical across all three implementation plans, because the same system is built regardless of rollout percentage. Fifteen components are costed individually; twelve are USD vendor-benchmarked and uplifted for India regulated-sector delivery and FX, three are priced directly in rupees.

Ref	Component	Rs Cr	Share	Phase
D-09	Trust Score orchestration core	1.878	19.0%	P4
D-06	CV damage assessment, 360° video and depth	1.002	10.1%	P2
D-08	Fraud and duplicate graph engine	1.002	10.1%	P3
D-12	Systems integration	1.002	10.1%	Cross
D-10	Production hardening	0.751	7.6%	P4
D-04	Synthetic and tampered image screening	0.626	6.3%	P2
D-13	Training data: labelled damage/synthetic set	0.600	6.1%	P2
D-07	Repair cost estimation engine	0.501	5.1%	P2
D-01	Document AI, OCR and field extraction	0.438	4.4%	P1
D-11	Claimant, surveyor and garage consoles	0.438	4.4%	P4
D-14	Model QC and UAT window	0.400	4.0%	Cross
D-05	Guided live-capture SDK, gallery disabled	0.376	3.8%	P1
D-02	Policy validation, RAG on actual wording	0.313	3.2%	P1
D-03	EXIF and metadata forensics engine	0.313	3.2%	P1
D-15	Security, DPDP and SOC2 setup	0.250	2.5%	Cross
W-25	TOTAL BUILD COST	9.890	100%	All

Table D. USD lines are uplifted by the India regulated-sector cost factor (A-10, 1.314x) and converted at spot FX (A-08, Rs 95.3/USD); D-13, D-14 and D-15 are priced directly in rupees.



Reading the stack


The three Capture Integrity Gate components (D-03, D-04, D-05) plus the training data line (D-13) together total Rs 1.91 Cr — the single largest functional block once summed, and by design: it is the platform's evidence layer and its primary competitive moat. Production hardening (D-10) and the security/DPDP setup (D-15) together total Rs 1.00 Cr of governance and auditability spend, carried as a cost line rather than assumed away.

NOTE 13 ANNUAL OPERATING COST SCHEDULE

Six lines make up the annual run cost: two scale with claim volume, four are fixed regardless of how many claims the platform processes.

Ref	Cost line	Conservative	Base	Aggressive	Type
W-26	GPU compute, peak-sized	0.63	1.88	3.14	Variable
W-27	Storage	0.07	0.20	0.34	Variable
W-28	MLOps tooling	0.80	0.80	0.80	Fixed
W-29	Operations team	0.64	0.64	0.64	Fixed
W-30	Security and compliance monitoring	0.25	0.25	0.25	Fixed
W-31	Legal and regulatory counsel	0.12	0.12	0.12	Fixed
W-32	TOTAL ANNUAL RUN COST	2.51	3.90	5.29	—

Table E (Base plan run stack, Rs 3.90 Cr total — live from Sheet 3 Part E, W-26 through W-32. The Cost Architecture sheet's Rs 3.68 Cr figure is a stale, hand-typed snapshot that was never updated after the Sheet 3 model was revised and should not be quoted).



The ceiling

	Rs Cr / year	Share of run cost
Variable lines (GPU + storage)	2.086	53.5%
Fixed lines (tooling + ops + security + legal)	1.815	46.5%

No cost-optimisation programme, however well executed, can take annual run cost below approximately Rs 1.81 Cr at this operating design — that is the fixed floor of tooling, headcount, security monitoring and legal counsel. Every optimisation decision in Note 14 targets the variable half only.

NOTE 14 AI COST ARCHITECTURE — THE NINE OPTIMISATION DECISIONS

Nine design decisions were taken specifically to hold GPU compute — the single largest variable line at 48.3% of run cost — as low as possible. Together they attack the Rs 2.09 Cr variable half (GPU plus storage); the remaining Rs 1.81 Cr fixed floor is untouched by any of them. This note states that ceiling explicitly rather than leaving it to be discovered under questioning: cost optimisation is not what makes the investment case work. Even eliminating every rupee of compute moves payback by roughly 0.1 month. The nine decisions exist to prove the run-cost line is bounded and will not run away at scale, not to manufacture savings.

#	Decision	Cost Line Attacked	Status
1	Directed intake form, not free-text description	GPU / inference volume	Designed
2	Policy wording parsed once, offline, into a clause table	GPU, per-claim policy processing	Designed
3	OCR before any multimodal call	GPU	Designed
4	Video de-duplication and damage-frame selection	GPU (largest single lever)	Designed, ratio not yet measured
5	Fraud scoring queued in 10-minute clusters	GPU, twice over, plus detection quality	Designed
6	Peak flattening through queueing	GPU peak-to-average multiplier	Designed, multiplier not yet measured
7	Context caching, system prompt and shared context	Any API-path inference cost	Available on the API path
8	Model routing, economy tier for triage	Any API-path inference cost	Available on the API path
9	Deterministic-first on the green lane	GPU	Designed, already in the lane mix

Table B, Cost Architecture sheet. "Available" decisions are priced but deliberately left unbanked in the headline figures — genuine reserve headroom, not counted benefit.

14.1 Directed intake form, rather than free-text description

Sits at the FNOL layer, before the Capture Integrity Gate. A free-text claim description has to be parsed by a model before anything useful can be extracted from it; a structured form with defined fields removes that parsing step entirely, because there is no prose to interpret. This was taken as the default design, not an optimisation retrofitted later, because it also yields a second benefit at zero marginal cost: completion time, field-revision pattern, and paste-versus-type behaviour are all early fraud signals the free-text alternative would not provide. It sets the LOW case in the token-range test at Note 15: 4,962 tokens per claim, against 88,720 in the HIGH case.

14.2 Policy wording parsed once, offline, into a clause table

Lives inside the Policy Validation RAG engine (build cost Rs 0.313 Cr). Policy wording is byte-identical across every claim on a given product, so parsing the full wording on every single claim is pure waste — the decision was to parse each product's wording once, offline, into a structured clause table, so a claim-time lookup replaces a claim-time inference call. This converts a per-claim cost into a one-time per-product cost, and is worth 40,000 tokens per claim at the top of the token range tested at Note 15 — the single largest line item in the HIGH case.

14.3 OCR before any multimodal call

Sits ahead of the Document AI engine. Reading text off a scanned document using a vision-capable multimodal model is the expensive way to extract a few hundred characters; deterministic OCR is cheap, fast, and produces the same extracted fields. The decision was to run OCR first and only pass the extracted structured fields downstream, never the raw document image, to any multimodal model call. No isolated rupee figure exists for this decision alone, since it is a routing choice rather than a priced component, but it keeps the Document AI engine (D-01) deterministic and cheap by construction.

14.4 Video de-duplication and damage-frame selection

Sits inside the CV Damage Assessment engine (build cost Rs 1.002 Cr, the largest engine after Trust Score orchestration). Video tokenises at 263 tokens per second — a bare 40-second claim video is already 10,520 tokens before any other line is added, making it the dominant cost driver in every scenario tested. The decision was to apply perceptual hashing to drop near-duplicate frames and a cheap classifier to retain only frames that actually show damage, before the expensive damage-assessment model ever runs. This is flagged explicitly as the highest-value lever on the cost-architecture sheet — and also as the one honest gap: the achievable reduction ratio has not yet been measured, so no rupee saving is claimed for it. It is a designed mechanism awaiting pilot instrumentation, not a banked number.

14.5 Fraud scoring queued in 10-minute clusters

Sits at the Fraud and Duplicate Graph engine (build cost Rs 1.002 Cr). Two independent effects justify batching rather than scoring each claim the instant it arrives. First, every vendor rate card publishes a 50% discount for asynchronous batch processing over synchronous calls — a pure cost effect. Second, and more importantly, a fraud ring spread across several claims submitted close together cannot be detected by a model scoring one claim in isolation; batching claims into 10-minute clusters lets the graph engine see connections a one-at-a-time design would miss entirely. The batching discount is priced but not applied to the headline figures — it is treated as unused headroom, while the detection-quality improvement it enables is already folded into the modelled fraud-detection uplift. The Green lane is deliberately kept synchronous throughout, so auto-settlement is never delayed by this batching.

14.6 Peak flattening through queueing

Applies at the infrastructure layer, where it sizes the GPU compute provisioning line directly. GPU capacity is currently provisioned for a peak-to-average multiplier of 2.5x, with a further 1.2x safety margin on top — meaning the system is built to handle roughly three times average load at all times, because claims cluster around specific hours and events rather than arriving evenly. The decision was to route non-urgent processing (Amber and Red lane work that does not need instant turnaround) through a queue that smooths this peak rather than provisioning hardware for the peak itself. The achievable reduction in the peak multiplier has not yet been measured against live claim-arrival data, so — as with decision 4 — this is a designed mechanism with a real cost lever behind it, not a number already banked into the run-cost figure.

14.7 Context caching of the system prompt and shared context

Applies across all five engines wherever an API path is used. The system prompt and the product policy context are identical on every single claim processed, and all three inference vendors (Google, OpenAI, Anthropic) publish roughly a 90% discount on cached, re-used input tokens compared with fresh input. This is priced and verified against vendor rate cards dated 12 August 2026, but is deliberately not applied to any headline cost or benefit figure in this model — it exists as quantified, available headroom, not a claimed saving, and only applies where an API path (rather than self-hosted inference) is actually chosen.

14.8 Model routing — economy tier for triage

Applies at the Trust Score orchestration layer, which decides which model tier serves each engine call. The verified blended economy-tier input rate is USD 0.50 per million tokens, against USD 1.83 for the blended mid-tier — a more than 3.5x difference for the same task. The decision was that triage and simple classification work (deciding which lane a claim likely belongs in before the full engine suite runs) does not need the most capable, most expensive model tier; only genuinely difficult judgement calls do. Applying this routing would drop the API-path HIGH case from Rs 19.41 to Rs 2.67 per claim — again priced but treated as unbanked headroom, available if and when an API delivery path is adopted for any part of the pipeline.

14.9 Deterministic-first on the Green lane

Applies at the Trust Score / routing layer, after all five engines have fused their outputs into a Trust Score. The cheapest inference call is the one that is never made at all: 65% of claims are designed to clear on rule-based checks and Capture Integrity verification alone, inside the IRDAI Rs 50,000 surveyor-exemption corridor, without any model call being invoked. This decision is already fully reflected in the stated run-cost and benefit figures throughout this document — it


is the structural reason the Rs 3.90 Cr annual run stack is not considerably larger, since well over half of claim volume never reaches the more expensive engine calls in the first place.

NOTE 15 INFERENCE COST SENSITIVITY AND BREAK-EVEN

Actual per-claim token consumption has not been measured — no pilot has run yet. Rather than publish a single unverifiable point estimate, three cases (LOW, CENTRAL, HIGH) are tested across an 18-fold range, and the finding is that the range does not matter to the investment decision.


15.1 Token composition, three cases

Component	LOW	CENTRAL	HIGH	Basis
Video seconds analysed	10	40	120	40s is Engine 02's stated derivation
Photo tiles	4	8	20	258 tokens/tile, count assumed
Policy wording tokens	0	8,000	40,000	Zero if parsed offline (decision 2)
Claim text / document tokens	300	1,500	4,000	300 = structured form (decision 1)
System prompt tokens	1,000	3,000	8,000	Cacheable at ~90% off (decision 7)
Output tokens	300	1,200	4,000	300 = structured JSON decision
TOTAL INPUT TOKENS	4,962	25,084	88,720	Sum of the above




15.2 Cost and payback impact

Case	Mid-tier Rs/claim	Mid-tier Rs Cr/yr	Economy+batch Rs/claim	Payback
LOW	1.16	0.017	0.16	3.074 mo
CENTRAL	5.56	0.082	0.76	3.079 mo
HIGH	19.41	0.286	2.67	3.096 mo
No API cost at all	—	—	—	3.073 mo



Input token volume varies eighteen-fold between the LOW and HIGH cases, yet payback moves by only 0.02 months. An order-of-magnitude error in the token estimate would not change the investment decision — this insensitivity, not the point estimate itself, is the finding worth presenting to a panel.

15.3 Breakeven — how large would inference have to become before it mattered



A ClaimPulse claim carries forty seconds of video. Inference consumption would have to rise roughly fifteen-fold — to about ten minutes of video per claim — before it consumed even 1% of net annual benefit, and to about ninety-five minutes before it consumed 10%. This is the complete answer to whether AI running costs threaten the case, and it does not depend on the token estimate being exactly right.

NOTE 16 WHAT OPTIMISATION IS WORTH

Having shown the run-cost line is bounded (Note 15), this note prices the ceiling of what the nine decisions in Note 14 could ever be worth, to make plain that cost optimisation is a credibility exercise, not a value driver, in this investment case.

Scenario	Run cost Rs Cr	Net annual Rs Cr	Payback, months	Basis
As modelled today	3.90	43.31	2.74	Live from Sheet 3, H-08/H-09/H-12
GPU line reduced 80%	2.39	44.82	2.65	Illustrative, not measured
GPU line removed entirely	2.02	45.19	2.63	Upper bound, not achievable
Fixed floor	1.81	—	—	Cannot go below this at any optimisation level

The middle two rows are illustrative, not measured reductions. Even in the impossible upper-bound case of removing all GPU compute cost entirely, payback improves by 0.11 months — from 2.74 to 2.63. Cost optimisation is presented here to prove the run-cost line is defensible and will not run away with scale, not because it moves the investment case in any material way. Any presentation of these rows as achieved or committed savings would misrepresent the model.

NOTE 17 GENAI BUILD VERSUS BUY


Bajaj needs conventional-ML compute for rule-based fraud detection regardless of any GenAI decision, so that baseline is excluded from this comparison. This note isolates only the workload where a genuine API alternative exists — document extraction, policy reasoning, and explanation drafting — and compares five-year self-hosted compute cost against five-year API cost for that workload alone.

17.1 Scope of the decision

Ref	Component	Share of GPU Hours	Basis
O-01	Baseline conventional ML, needed regardless	25%	Excluded — required either way
O-02	Moat, self-hosted regardless	30%	Excluded — no buy option exists
O-03	In-scope: GenAI-substitutable workload	45%	Document/policy/explanation work

17.2 The five-year verdict

Ref	Measure	Value	What It Means
T-07	Five-year NPV, BUY	Rs 0.313 Cr	Total incremental cost of the API route
T-08	Five-year NPV, BUILD	Rs 4.055 Cr	Compute only — a floor, since it excludes serving/engineering overhead
T-09	NPV advantage to BUYING	Rs 3.742 Cr	Positive means buying is cheaper
T-11	Provisioned GPU-seconds/claim implied by the model	648	Current assumption
T-12	Provisioned GPU-seconds/claim at breakeven	77	Below this, self-hosting wins
T-13	Gap over breakeven	8.4x	How far current design sits from the flip point




Both options scale linearly with claim volume, so there is no crossover volume — doubling claims doubles both sides and leaves the ranking unchanged. The decision is set entirely by GPU-seconds required per claim for the in-scope workload, which is currently a Tier 4, unmeasured assumption (W-09 in the sensitivity register). The 90-day pilot is designed to measure this directly.

17.3 Why self-host anyway, in part

Reason	Detail
DPDP and data residency	Claim media carries faces, number plates and location. Both OpenAI and Anthropic charge a 10% premium for region-pinned inference, and cross-border processing of claimant personal data is a live, unsettled compliance question.
The moat cannot be rented	The synthetic/tampered image classifier is trained on Bajaj's own claim media. No general-purpose API performs diffusion/GAN artefact detection tuned to Indian motor claims — outsourcing this would be outsourcing the differentiator.
Vendor and price risk	30% annual inference deflation is a tailwind for buying, but it is a vendor-controlled price. A single external supplier setting the marginal cost of the core product is a strategic exposure worth pricing in.
<i>The honest conclusion: the defensible architecture is split, not either/or. Self-host the Capture Integrity Gate and the fraud graph — the moat and the DPDP-sensitive media. Buy the document, policy and explanation work — commodity language processing on text already de-identified by OCR. Every cost this note omits (integration, serving infrastructure, hardware refresh, latency, switching cost) understates BUILD more than it understates BUY, so correcting for them would widen the gap in favour of buying, not close it.</i>	

NOTE 18 NET RESULT AND PAYBACK

Ref	Metric	Conservative	Base	Aggressive	Unit
W-33	Gross annual benefit	15.36	47.21	80.57	Rs Cr
W-34	Annual run cost	2.51	3.90	5.29	Rs Cr
W-35	NET ANNUAL BENEFIT	12.85	43.31	75.28	Rs Cr
W-36	Net monthly benefit	1.07	3.61	6.27	Rs Cr
W-37	Build cost, one-off	9.89	9.89	9.89	Rs Cr
W-38	Payback, steady-state basis	9.24	2.74	1.58	months




Steady-state payback divides build cost by net monthly benefit and is the standard convention for the measure. It excludes the ten-month build window during which no benefit accrues — the project-kickoff figure that includes that window is presented in Note 27 and is the number this proposal's headline quotes.

NOTE 19 INSURANCE P&L BRIDGE — COMBINED RATIO IMPACT

Combined ratio equals loss ratio plus expense ratio. Fraud avoidance reduces incurred claims and moves the loss ratio; labour savings sit in claims-handling opex and move the expense ratio. Both are struck against Motor OD net earned premium, and separately restated against Bajaj's total book, because the two denominators must never be conflated.

Ref	Metric	Conservative	Base	Aggressive	Unit
W-39	Motor OD net earned premium	1,654.67	1,654.67	1,654.67	Rs Cr
W-40	Rs per 1pp of Motor OD ratio	16.55	16.55	16.55	Rs Cr
W-41	Loss ratio impact (all three fraud lines)	0.47	1.42	2.37	pp
W-42	Expense ratio impact (labour, net of friction and run cost)	0.29	1.16	2.13	pp
W-43	COMBINED RATIO IMPACT — Motor OD basis	0.76	2.58	4.49	pp
W-44	Bajaj total net earned premium, estimated	10,205.95	10,205.95	10,205.95	Rs Cr
W-45	Rs per 1pp of group ratio	102.06	102.06	102.06	Rs Cr
W-46	COMBINED RATIO IMPACT — Group basis	0.12	0.42	0.73	pp
W-47	Distribution income (renewal), outside the combined ratio	0.19	0.58	0.96	Rs Cr



Never subtract the Motor OD figure from a group combined ratio — the two rows above are the same rupees measured against two different denominators, not two different amounts of benefit. Renewal retention (W-47) is distribution commission, not an underwriting item, and is excluded from both ratios.

NOTE 20 DOWNSIDE AND CORRELATED STRESS TESTS


20.1 Single-lever downside tests

Ref	Test line	Conservative	Base	Aggressive	Unit
W-48	Labour benefit at the in-house touch cost (Rs 83)	2.63	7.88	13.14	Rs Cr
W-49	Fraud avoided if detection lands at 82%	4.09	12.28	20.46	Rs Cr
W-50	Less: friction cost at the in-house rate	-0.07	-0.20	-0.33	Rs Cr
W-51	Downside net annual benefit	4.14	16.06	27.98	Rs Cr
W-52	Floor net annual benefit (labour only, in-house, zero fraud)	0.05	3.78	7.52	Rs Cr
W-53	Payback — as modelled	9.24	2.74	1.58	months
W-54	Payback — realistic downside	28.63	7.39	4.24	months
W-55	Payback — absolute floor	2,286.0	31.36	15.79	months

The realistic downside (W-51) combines in-house labour cost, fraud detection landing at 82% instead of 90%, and zero synthetic-media or renewal benefit — a genuinely conservative composite, not a single-variable stress. The absolute floor (W-52) removes fraud benefit entirely, leaving only labour savings at in-house rates; at Conservative rollout this floor is barely positive, which is disclosed rather than hidden.

20.2 Correlated stress test — several assumptions missing at once (Base plan)

Lever / output	Base	A: Operating reality	B: Fraud disappoints	C: Adoption drag	D: All three together
Baseline cost per touch, Rs	250	83	250	250	83
Baseline touches per claim	7	5	7	7	5
Fraud detection rate, target	90%	90%	82%	90%	82%
Rollout of the addressable book	60%	60%	60%	35%	35%
Green lane share	65%	65%	65%	50%	50%
GROSS ANNUAL BENEFIT, Rs Cr	47.21	29.02	31.15	26.59	7.24
Annual run cost, Rs Cr	3.90	3.90	3.90	3.03	3.03
NET ANNUAL BENEFIT, Rs Cr	43.31	25.11	27.25	23.56	4.21
5-YEAR NPV at 12%, Rs Cr	145.94	79.55	87.33	74.36	3.76
Payback from project kickoff, months	16.84	22.36	21.73	22.49	47.66
INVESTOR GATE — 5-yr NPV still positive?	PASS	PASS	PASS	PASS	PASS




Scenario A moves operating assumptions to in-house handling (Rs 83/touch, five baseline touches). B lands fraud detection at 82%, halves the leakage pool, and cuts synthetic-media incidence to a quarter. C holds rollout at 35%, Green lane at 50%, and doubles live-capture friction. D applies all three simultaneously. The build is Rs 9.89 Cr; even the fully compounded downside (D) repays it and remains NPV-positive at Rs 3.76 Cr.

NOTE 21 STAKEHOLDER VALUE SPLIT

Every rupee of net benefit is attributed to the entity that books it, so the split reconciles exactly to the group total with no unallocated residual.

Ref	Metric	Conservative	Base	Aggressive	Unit
W-56	Bajaj General claims ops — expense-ratio value	7.34	23.15	40.47	Rs Cr
W-57	Bajaj General underwriting — loss-ratio value	7.83	23.49	39.14	Rs Cr
W-58	BFDL distribution — renewal value	0.19	0.58	0.96	Rs Cr
W-59	Less: platform run cost	-2.51	-3.90	-5.29	Rs Cr
W-60	CHECK — split less W-35, must be zero	0.00	0.00	0.00	Rs Cr



Operational and stakeholder-facing metrics

Ref	Metric	Conservative	Base	Aggressive	Unit
W-62	Claimant-days returned per year	388,968	1,166,904	1,944,839	days
W-63	Claims auto-settled with zero human touch	35,660	106,980	178,300	claims
W-64	Claims still reviewed by a human	19,202	57,605	96,008	claims
W-66	Manual touches avoided per year	316,551	949,652	1,582,754	touches
W-68	FTE capacity released	58.6	175.9	293.1	FTE
W-69	Adjuster throughput multiple	5.69x	5.69x	5.69x	x
W-72	Surveyor visits avoided per year	24,688	74,063	123,438	visits
W-78	Honest claimants downgraded green to amber	4,389	13,167	21,945	claims
W-79	Cost of that friction, per year	0.24	0.59	0.79	Rs Cr

FTE capacity released (W-68) is redeployed to Amber/Red judgement work, not eliminated — the adjuster throughput multiple (W-69) shows the same team clears 5.69 times as many claims. Honest claimants downgraded from Green to Amber (W-78) is disclosed as a named cost of the hard live-capture rule, not folded silently into the benefit lines.

NOTE 22 BREAK-EVEN ANALYSIS

What would have to be true for this to fail, tested on labour benefit alone, with every fraud, capture-gate, synthetic-media and renewal line set to zero.

Ref	Metric	Conservative	Base	Aggressive	Unit
W-90	Break-even manual touches per claim today	3.31	2.18	1.76	touches
W-91	Cushion on touches per claim	53%	69%	75%	ratio
W-92	Break-even baseline cost per touch	118.07	77.79	62.70	Rs
W-93	Cushion on cost per touch	53%	69%	75%	ratio
W-94	Break-even rollout of the addressable book	34.0%	32.4%	31.0%	ratio

Break-even here means net annual benefit falls to zero with the fraud, capture-gate, synthetic-media and renewal lines all set to zero — build cost is excluded, since it is a one-off cost addressed separately by payback. At Base plan, the 7-touch baseline assumption could fall to 2.18 touches (a 69% cushion) and the labour engine alone would still cover the full annual run cost, before a single rupee of fraud benefit is counted.

NOTE 23 CLAIM FREQUENCY RECONCILIATION

An integrity check: does the modelled claim count, built bottom-up from claims value divided by average claim size, reconcile against an independent estimate of policy count and claim frequency?


Ref	Metric	Value	Unit
W-96	Motor OD gross direct premium	3,317.30	Rs Cr
W-97	Average Motor OD premium per policy	7,000	Rs
W-98	Implied Motor OD policies in force	4,739,000	policies
W-99	Modelled annual claim count, full book	274,307	claims
W-100	IMPLIED CLAIM FREQUENCY	5.79%	ratio
W-101	Industry reference band, motor own damage	5% – 15%	ratio
W-102	CHECK — frequency inside industry band	PASS	text

Average claim size (B-01) and average premium (B-21) are independent Tier 4 inputs. Taken together they imply a claim frequency of 5.79% — inside the 5%–15% industry norm for motor own-damage, sitting near the lower bound because the claim count is built on net incurred claims rather than a gross-basis reading, which is the cross-check that makes both underlying inputs defensible.

NOTE 24 ALTERNATIVE NET EARNED PREMIUM BASES

The model runs throughout on the whole-book FY2026 NEP/GDPI ratio (49.88%) as its default basis. Two alternative bases are carried for transparency, since Motor OD retention runs above the whole-book blend, meaning the default basis is the conservative choice.


Basis	NEP Ratio	Motor OD NEP, Rs Cr	Claims Pool, Rs Cr	Net Annual Benefit, Rs Cr	Payback, Kickoff	Comb. Ratio, pp
Whole-book FY2026 (model default)	49.88%	1,654.67	1,234.38	43.31	16.84 mo	2.58
Whole-book Q1 FY2027	46.10%	1,529.28	1,140.84	39.89	17.47 mo	2.57
Motor OD retention cohort	68.85%	2,283.96	1,703.83	60.47	14.82 mo	2.61



Motor OD retention (the third basis) is the more representative measure for a Motor-OD-only book, and the case is materially stronger on it — Rs 60.47 Cr net annual benefit against the Rs 43.31 Cr headline. The gap between the three rows is the size of the conservatism built into the number this proposal actually quotes.

NOTE 25 THREE-YEAR FORECAST BY IMPLEMENTATION PLAN

Year 0 is the ten-month build with no benefit. Benefit ramps through Year 1 (45% realisation) and Year 2 (92%), reaching steady state in Year 3 (100%). Run cost is charged in full from Year 1 — a deliberately conservative convention, since it assumes full operating cost from go-live even while benefit is still ramping.



25.1 Base plan detail — 60% rollout

Line	Year 0	Year 1	Year 2	Year 3	3-yr total
Claims processed	—	74,063	164,591	194,468	433,122
Gross benefit, Rs Cr	—	21.24	47.21	55.78	124.24
Run cost, Rs Cr	—	3.90	4.24	4.61	12.75
Build cost, Rs Cr	9.89	0	0	0	9.89
NET BENEFIT, Rs Cr	-9.89	17.34	42.97	51.17	111.49
Risk-adjusted net benefit, Rs Cr	-9.89	13.22	32.77	39.02	85.01
Cumulative net position, Rs Cr	-9.89	7.45	50.43	101.60	—

Conservative and Aggressive plans follow the identical structure at 20% and 100% rollout respectively; only claims volume, gross benefit and run cost scale with rollout. Their year-by-year detail is omitted here for brevity — the trajectory chart above and the summary table below carry their full results.


25.2 Forecast summary — all three plans

Metric	Conservative	Base	Aggressive	Unit
Build cost (Year 0)	9.89	9.89	9.89	Rs Cr
3-year net benefit	32.22	111.49	194.73	Rs Cr
3-year net benefit, risk-adjusted	24.57	85.01	148.48	Rs Cr
3-year NPV at 12%	14.92	76.28	140.71	Rs Cr
PAYBACK FROM PROJECT KICKOFF	27.21	16.84	13.83	months

Payback from project kickoff includes the ten-month build window before any benefit accrues. This is the project-basis figure this proposal's headline quotes; the shorter steady-state figure in Note 18 excludes the build window.

NOTE 26 UNIT ECONOMICS PER CLAIM

Ref	Metric	Conservative	Base	Aggressive	Unit
UE-01	Cost to serve per claim — today	1,750	1,750	1,750	Rs
UE-02	Blended manual touches after ClaimPulse	1.23	1.23	1.23	touches
UE-03	Cost to serve per claim — after	369	307.50	246	Rs
UE-04	Labour saving per claim	1,381	1,442.50	1,504	Rs
UE-05	Run cost per claim	457.49	237.00	192.90	Rs
UE-06	Fraud/synthetic exposure avoided per claim	1,426.97	1,426.97	1,426.97	Rs
UE-07	NET BENEFIT PER CLAIM	2,342.28	2,631.48	2,744.28	Rs
UE-08	Cost to serve as % of claim value — today	3.89%	3.89%	3.89%	ratio
UE-09	Cost to serve as % of claim value — after	1.84%	1.21%	0.98%	ratio
UE-10	Claim TAT, on-platform	2.71	2.71	2.71	days
UE-11	Break-even rollout, labour only	34.0%	32.4%	31.0%	ratio




Cost to serve per claim falls from Rs 1,750 to Rs 307.50 at the Base plan — an 82.4% reduction — even after fully loading the platform's own run cost into the after-figure (UE-03 includes UE-05). Net benefit per claim (UE-07) is the single truest per-unit economics measure in the model, since it nets every cost and every benefit line against actual claims processed.

NOTE 27 PROJECT DCF, NPV AND IRR

Five-year discounted cash flow, using the same benefit-realisation ramp as Note 25 (45% / 92% / 100% / 100% / 100%) and the Year-3 market-growth factor held flat for Years 4–5. Discounted at the 12% WACC hurdle rate.


Year	Conservative	Base	Aggressive	Discount Factor
Year 0 — 10-month build	-9.89	-9.89	-9.89	1.000
Year 1 — benefit ramping	4.40	17.34	30.96	0.893
Year 2	12.63	42.97	74.82	0.797
Year 3	15.18	51.17	88.95	0.712
Year 4 — steady state	16.50	55.63	96.68	0.636
Year 5 — steady state	17.94	60.47	105.10	0.567
5-year cumulative, undiscounted	56.77	217.69	386.62	—
5-YEAR NPV AT 12%	35.59	145.94	261.79	—
5-year IRR	91.0%	255.1%	406.6%	—



A small, fixed one-off build (Rs 9.89 Cr) set against a very large recurring claims pool produces a high IRR by construction. This proposal leads with payback and combined-ratio impact — the measures that scale sensibly with the size of the book — and carries IRR here in the annexure for completeness rather than as a headline figure.

NOTE 28 ROIC AND ECONOMIC VALUE ADDED

Ref	Metric	Conservative	Base	Aggressive	Unit
RE-01	Net annual benefit, steady state	12.85	43.31	75.28	Rs Cr
RE-02	Net annual benefit, realistic downside	4.14	16.06	27.98	Rs Cr
RE-03	Capital deployed — build only	9.89	9.89	9.89	Rs Cr
RE-04	Capital deployed — fully loaded (+15% contingency)	11.37	11.37	11.37	Rs Cr
RE-05	Hurdle rate (WACC)	12%	12%	12%	ratio
RE-06	ROIC — headline, on fully loaded capital	113.0%	380.8%	661.8%	ratio
RE-07	ROIC — realistic downside	36.4%	141.2%	246.0%	ratio
RE-08	EVA at 12% hurdle — headline	11.49	41.95	73.91	Rs Cr
RE-09	EVA at 12% hurdle — downside	2.78	14.70	26.62	Rs Cr
RE-10	Investor gate: does downside ROIC clear WACC?	PASS	PASS	PASS	test

ROIC — Even the Downside Case Clears the WACC Hurdle by a Wide Margin

ROIC is high on this project because a comparatively small software build (Rs 9.89–11.37 Cr) sits against a Rs 1,234 Cr annual claims pool — the ratio is a function of that scale mismatch, not an unusually favourable return per rupee. This proposal leads with payback and combined-ratio impact for that reason; ROIC and EVA are carried here for completeness and are always quoted on the realistic-downside basis (RE-07), not the headline, when the figure comes up.

NOTE 29 CAPITAL ALLOCATION — BUILD VERSUS BUY VERSUS WAIT


Alternative	3-yr TCO, Rs Cr	Time to Value	Strategic Control	Data Moat
Status quo: OTS plus manual handling	0	Immediate	High	None
Buy point solutions (OCR, fraud rules, or photo estimate)	Excluded on capability	6–12 months	Low	None
Buy a fused vendor platform	Excluded on capability	Live elsewhere	None	None
BUILD ClaimPulse — Conservative	18.09	10–11 months	Full	High
BUILD ClaimPulse — Base	22.64	10–11 months	Full	High
BUILD ClaimPulse — Aggressive	27.19	10–11 months	Full	High
Wait for motor pricing to correct	0	3–6 months	N/A	N/A
Buy the commodity AI layer inside the build	0.31	Immediate	Full	High

Point solutions and fused vendor platforms are excluded on capability, not on price — none fuses document, damage, policy, graph and capture evidence into a single routing decision, and a fused platform still screens no synthetic damage media, which every other engine depends on as its evidence layer. "Buy the commodity AI layer inside the build" is the Note 17 decision, already adopted within the Build option rather than a genuine standalone alternative.

NOTE 30 THE COST OF WAITING

The alternative to investing is not free. This note prices what continuing to operate on the current process costs, independent of whether ClaimPulse is ever built.

Ref	Metric	Value	Unit
CW-01	Motor OD loss ratio deterioration, Q1 FY27 vs Q1 FY26	0.032 pp	pp
CW-02	Annual rupee cost of that deterioration	52.95	Rs Cr
CW-03	Monthly opportunity cost of delay, Base plan	3.61	Rs Cr
CW-04	Cost of a six-month delay, Base plan	21.66	Rs Cr



Every month the decision is deferred forgoes one month of Rs 3.61 Cr net benefit at the Base plan. A six-month delay in approving this proposal costs approximately Rs 21.66 Cr in forgone benefit alone, before any further loss-ratio deterioration on the existing manual process is counted.

NOTE 31 STRATEGIC MOAT AND EARLY WARNING SYSTEM

The investment case should be judged against the system's breakpoints, not against a forecast. ClaimPulse's resilience does not come from predicting which future arrives — it comes from four structural mechanisms pulling in different directions at once: falling AI cost, rising digital evidence availability, rising complexity of verifying that evidence, and a governance ceiling on what may be automated. No single trend reversing breaks the case.

31.1 Strategic response, by scenario

If this materialises	The response is	Watch for
Experience parity commoditises speed	Invest ahead of the friction frontier, but test the marginal-value curve first — beyond a threshold, further TAT reduction may cost more than it returns	Competitor-advertised TAT; Bima Sugam TAT disclosure
Controlled autonomy tightens	Build governance as a production requirement, not a compliance layer — governance is what buys automation permission	How the IRDAI AI framework treats adverse decisions
Adversarial claims accelerate	Invest in Capture Integrity and synthetic-media detection — a risk-mitigation investment, since the labour engine already carries the case without fraud economics	Synthetic-media flag rate; false-negative fraud cases
AI components commoditise	Shift investment from AI components to data, orchestration, ecosystem and governance — buy the commodity layers, build the scarce ones (Note 17)	Vendor pricing; open-weight performance parity
A competitor reaches parity	Redirect investment toward owning the orchestration layer and the proprietary data feedback loop, not the individual AI components	Orchestration-layer build or RFP announcements, internal or competitor

31.2 Early warning — what to monitor every quarter

Signal	Observable indicators	Data owner
Evidence integrity	Gallery vs live-capture share; first-pass image acceptance; EXIF availability; reinspection rate; synthetic-media flag rate	ClaimPulse telemetry
Adversarial AI	Suspicious-media prevalence; false-negative fraud cases; fraud-ring clustering density; classifier drift	SIU and ClaimPulse telemetry
Conditional autonomy	Human override rate; adverse-decision appeal rate; Amber-to-Red migration rate	ClaimPulse telemetry and Ombudsman data
AI commoditisation	Model cost per token; open-weight performance parity; competitor deployment announcements	Technology function
Governance tightening	IRDAI AI framework publications; DPDP enforcement actions; explainability challenge rate	Compliance and Legal

Most of these indicators are ClaimPulse's own telemetry, which begins the day the system goes live — this is the argument for the 90-day pilot in Note 2: it converts these from unmeasurable to monitored from week one, not from year two.


31.3 The five interactions that determine the case

Variable	Why it matters more than any single-technology claim
The adversarial gap	Not fraud capability in absolute terms, but detection minus generation capability. This single variable determines whether fraud economics contribute or evaporate.
The automation frontier	Not what AI can do, but what AI is permitted to decide — set by the lower of technical capability and trust/governance permission.
The evidence integrity function	Not how much evidence arrives, but what share is trustworthy enough to auto-decide on. Digitalisation times evidence integrity gives automation eligibility, not digital volume alone.
The orchestration ownership question	Not whether AI orchestrates claims, but who owns the orchestration layer. If a platform, OEM or TPA builds it first, the moat weakens regardless of component quality.

Variable	Why it matters more than any single-technology claim
The evidence usability gap	Not how much evidence technology makes available, but what proportion is legally and operationally usable once consent, provenance and API fragmentation are accounted for.

NOTE 32 IMPLEMENTATION ROADMAP

The build proceeds through four gated phases over ten months, matching the Year-0 build convention used throughout the financial model (Notes 25 and 27). Each phase has a named exit gate; no phase auto-triggers the next, and capital for the following phase is only released once its predecessor's gate is passed.



Phase	Delivers	Exit gate
Phase 0 (Months 0-3)	Foundations: data pipeline, systems integration, security and DPDP setup	Integration test pass; DPDP controls signed off by Legal
Phase 1 (Months 3-6)	Capture Integrity Gate, Document AI and OCR, Policy RAG	Synthetic-media detection meets its target rate on a held-out test set; policy-clause extraction accuracy validated against a sample
Phase 2 (Months 6-9)	CV damage assessment, fraud and duplicate graph, repair-cost engine	Damage-severity accuracy validated against surveyor ground truth; fraud graph shadow-scored against historical claims
Phase 3 (Months 9-10)	Trust Score orchestration, full parallel pipeline, production hardening	All five engines running concurrently in shadow mode; go-live readiness review

Every measurement identified as an open gap in Note 2 is scheduled inside this roadmap rather than left to be resolved informally: GPU-seconds per claim and token consumption are instrumented from Phase 0 onward; damage-frame reduction ratios and the achievable peak-to-average multiplier are measured through Phase 2; and the human exception-capacity ceiling (breakpoint H* in Note 8) is the first live-telemetry measurement taken once Phase 3 goes live.

NOTE 33 THE INVESTMENT ASK

Approve Rs 9.89 Cr in phased capital for ClaimPulse, released across the four gated phases in Note 32, against a base-case 5-year NPV of Rs 145.94 Cr at a 12% discount rate and a steady-state payback of 2.74 months.

33.1 What is being asked for

Item	Detail
Capital requested	Rs 9.89 Cr, one-off build cost (Note 12), identical across all three rollout plans
Structure	Released in four tranches against the phase gates in Note 32 — no tranche is released automatically
Recommended plan	Base — 60% rollout of the addressable Motor OD book
Headline return	5-year NPV Rs 145.94 Cr at 12%; steady-state payback 2.74 months; project-kickoff payback 16.84 months (Note 27)
Downside protection	Realistic-downside net annual benefit Rs 16.06 Cr, 7.39-month payback; the fully compounded correlated stress case remains NPV-positive at Rs 3.76 Cr (Note 20)
What the capital buys beyond the build	A proprietary fraud graph and synthetic-media classifier trained on Bajaj's own claim media — an asset that compounds with volume and cannot be purchased from a vendor (Note 7)

33.2 Why now

Three things are true simultaneously and none of them improve by waiting: Bajaj's Motor OD loss ratio has already deteriorated for two consecutive quarters (Note 3); the regulatory exemption this design relies on (sub-Rs 50,000 claims, no mandatory survey) is already in force rather than pending; and every month of delay forgoes a further month of the Rs 3.61 Cr Base-plan net monthly benefit — Rs 21.66 Cr over a six-month delay alone (Note 30).

33.3 What the committee is not being asked to accept on faith

Every headline figure in this ask is stress-tested, not merely modelled: the case survives a compounded scenario where rollout, fraud detection and adoption all disappoint at once (Note 20); it survives an 18-fold error in the AI token-consumption estimate (Note 15); and across seventeen independently researched scenarios spanning four research domains, exactly one combined technology shock turns monthly economics negative, and that boundary is stated openly rather than left to be discovered (Note 8). The six open measurement gaps that remain (Note 2) are scheduled inside the roadmap's own phase gates, not deferred indefinitely.